

Assignment 3 Tutorial

MSBD5009/COMP5112 - Parallel Programming

CUDA Programming

9 Apr 2022

Tutorial Overview

- Problem Description
- Task Description
- Implementation Instruction

Problem Description

- Supermer Generation
 - a **Read** has multiple **K-mers**
 - a **K-mer** has one **minimizer** (value)
 - consecutive **K-mers** with the same minimizer value will generate a **Supermer**

<i>Read</i> =	CAAATTACTGCATA
(k-mer #1)	<u>CAAATTACT</u>
(k-mer #2)	<u>AAATTACTG</u>
(k-mer #3)	<u>AATTACTGC</u>
(super-mer #1)	<u>CAAATTACTG</u>
super-mer #1 is made up of k-mer #1 and #2, minimizer	
(k-mer #4)	<u>ATTACTGCA</u>
(super-mer #2)	<u>AATTACTGC</u>
super-mer #2 is made up of k-mer #3 only, minimizer	
(k-mer #5)	<u>TTACTGCAT</u>
(k-mer #6)	<u>TACTGCATA</u>
(super-mer #3)	<u>ATTACTGCATA</u>
super-mer #3 is made up of k-mer #4 #5 #6, minimizer	

Task Description

- Supermer Generation with CUDA
 - Data parallelism
 - General idea - two stages:
 - 1. Generate minimizers of all k-mers in a read.
 - Read sequence: ACTGACTG ...
 - Its kmers: kmer #1, kmer #2, kmer #3 ...
 - Its minimizer sequence: minimizer for kmer#1, minimizer for kmer #2 ...
 - 2. Find supermer beginning positions
 - A new supermer begins when the current minimizer \neq the last minimizer
 - Find consecutive minimizers with the same value

Task Description

- The given function is **GenerateSupermer_GPU**
 - `void GenerateSupermer_GPU(vector<string> &reads, int K, int P, vector<string> &all_supermers, int NUM_BLOCKS_PER_GRID, int NUM_THREADS_PER_BLOCK)`
 - Organize read data in CSR format,
 - Malloc arrays on both host and device,
 - Malloc the host arrays to device arrays,
 - Call the two kernel functions,
 - Malloc device data back to host,
 - Generate and save the supermers to the vector.

Implementation Instruction

- You need to finish two kernel functions in **gensupermer_cuda.cu**:
 - GPU_GenMinimizer
 - `__global__ void GPU_GenMinimizer(T_GPU_data data, int K_kmer, int P_minimizer)`
 - Input: reads in CSR format
 - Output: minimizer sequences
 - GPU_GenSKM
 - `__global__ void GPU_GenSKM(T_GPU_data data, int K_kmer, int P_minimizer)`
 - Input: minimizer sequences
 - Output: supermer offset sequences

Implementation Instructions

- Data Structures
 - Stage 1: **GPU_GenMinimizer**
 - Input: reads in CSR format
 - ***char *reads***: the read data
 - ***int *reads_offs***: the read data offsets
 - Output: minimizer sequences
 - ***unsigned int *minimizers***: the 2bits-compressed minimizers (an unsigned int sequence whose length is $n-k+1$. You can use the offset array ***reads_offs*** for this array)
 - 2bits-compression: 'A' -> 0b00, 'C' -> 0b01, 'G' -> 0b10, 'T' -> 0b11
 - So that a minimizer ($p \leq 16$) can be saved in a 32-bit unsigned int for better accessibility and easy comparison.
 - d_basemap is for easy type casting from char to unsigned int.
 - A sample usage:
 - `Unsigned int minimizer = 0;`
 - In a for-loop do: `minimizer = (minimizer << 2) | d_basemap[read[i]];`
 - (e.g., $p=6$, the minimizer AAACCT will be transformed to an integer 0b 00 00 00 01 01 11.)

Implementation Instructions

- Data Structures
 - Stage 1: **GPU_GenMinimizer**
 - Input: reads in CSR format
 - *char *reads*: the read data
 - *int *reads_offs*: the read data offsets
 - Output: minimizer sequences
 - *unsigned int *minimizers*: the 2bits-compressed minimizers (you can use the offset array *reads_offs* for this array)
 - Example (2 reads, k=8, p=4)
 - *reads: AAAACCTT CAAACCTTGG
 - *reads_offs: 0, 8, 18
 - *read_len: 8, 10
 - *minimizers: 0b00000000, -, -, -, -, -, -, -,
0b00000001, 0b00000001, 0b00000101, -, -, -, -, -, -, -

Implementation Instructions

- Data Structures
 - Stage 2: **GPU_GenSKM**
 - Input: minimizer sequences
 - *unsigned int *minimizers*: the 2bits-compressed minimizers (this array shares the same offsets with *reads)
 - *int *reads_offs*: the read data offsets (used also as the minimizer sequence offsets)
 - Output: supermer offset sequences
 - *int *supermer_offs*: the beginning position in read of each supermer, the last item should be $n-k+1$ as the ending indicator (this array shares the same offset array *reads_offs* with *reads)
 - Example (2 reads, $k=8$, $p=4$)
 - *reads: AAAACCTT CAAACCTTGG
 - *reads_offs: 0, 8, 18
 - *minimizers: 0, -, -, -, -, -, -, -, 1, 1, 5, -, -, -, -, -, -, -
 - *supermer_offs: 0, 1, -, -, -, -, -, -, 0, 2, 3, -, -, -, -, -, -, -

The last item in each offset sequence is $n - k + 1$ (the underlined item).

Thank you

MSBD5009/COMP5112 - Parallel Programming

A3: CUDA Programming